

main

August 29, 2024

1 Laboratorio 2

1.0.1 Luis Santos

1.0.2 Carné 20226

```
[ ]: import tensorflow as tf
from tensorflow.keras import layers, models
import numpy as np
import matplotlib.pyplot as plt
import shap
import os
from ipywidgets import interact, IntSlider

print("TensorFlow version:", tf.__version__)
print("GPU is", "available" if tf.config.list_physical_devices('GPU') else "NOT_A_
↪AVAILABLE")
```

TensorFlow version: 2.17.0
GPU is NOT AVAILABLE

```
[ ]: (train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.
↪cifar10.load_data()
train_images, test_images = train_images / 255.0, test_images / 255.0
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', '
↪horse', 'ship', 'truck']
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 6s
0us/step

```
[ ]: def create_model():
    model = models.Sequential([
        layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(64, (3, 3), activation='relu'),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(64, (3, 3), activation='relu'),
        layers.Flatten(),
```

```

        layers.Dense(64, activation='relu'),
        layers.Dense(10, activation='softmax')
    ])
    model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
metrics=['accuracy'])
    return model

```

```

[ ]: model_path = 'cifar10_cnn.keras'
if os.path.exists(model_path):
    print("Loading pre-trained model...")
    model = tf.keras.models.load_model(model_path)
else:
    print("No pre-trained model found. Training a new model...")
    model = create_model()
    history = model.fit(train_images, train_labels, epochs=10,
validation_data=(test_images, test_labels))
    model.save(model_path)
    print(f"Model saved to {model_path}")

```

No pre-trained model found. Training a new model...

/Users/lis/Library/Python/3.9/lib/python/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Epoch 1/10

1563/1563 15s 9ms/step -

accuracy: 0.3417 - loss: 1.7794 - val_accuracy: 0.5462 - val_loss: 1.2780

Epoch 2/10

1563/1563 15s 10ms/step -

accuracy: 0.5705 - loss: 1.2058 - val_accuracy: 0.6167 - val_loss: 1.0846

Epoch 3/10

1563/1563 16s 10ms/step -

accuracy: 0.6334 - loss: 1.0416 - val_accuracy: 0.6181 - val_loss: 1.0782

Epoch 4/10

1563/1563 15s 10ms/step -

accuracy: 0.6726 - loss: 0.9381 - val_accuracy: 0.6586 - val_loss: 0.9537

Epoch 5/10

1563/1563 15s 10ms/step -

accuracy: 0.7029 - loss: 0.8568 - val_accuracy: 0.6736 - val_loss: 0.9291

Epoch 6/10

1563/1563 15s 10ms/step -

accuracy: 0.7215 - loss: 0.7955 - val_accuracy: 0.6959 - val_loss: 0.8733

Epoch 7/10

1563/1563 14s 9ms/step -

accuracy: 0.7371 - loss: 0.7552 - val_accuracy: 0.6949 - val_loss: 0.8884

```

Epoch 8/10
1563/1563          15s 10ms/step -
accuracy: 0.7548 - loss: 0.7053 - val_accuracy: 0.6996 - val_loss: 0.8842
Epoch 9/10
1563/1563          15s 10ms/step -
accuracy: 0.7679 - loss: 0.6661 - val_accuracy: 0.6992 - val_loss: 0.8754
Epoch 10/10
1563/1563          15s 10ms/step -
accuracy: 0.7781 - loss: 0.6291 - val_accuracy: 0.7017 - val_loss: 0.9083
Model saved to cifar10_cnn.keras

```

```
[ ]: test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
print(f"Test accuracy: {test_acc:.2f}")
```

```

313/313 - 1s - 4ms/step - accuracy: 0.7017 - loss: 0.9083
Test accuracy: 0.70

```

```
[ ]: def test_model(model, test_images, test_labels, num_samples=5):
    indices = np.random.choice(test_images.shape[0], num_samples, replace=False)
    sample_images = test_images[indices]
    sample_labels = test_labels[indices]

    predictions = model.predict(sample_images)

    fig, axes = plt.subplots(1, num_samples, figsize=(15, 3))
    for i, ax in enumerate(axes):
        ax.imshow(sample_images[i])
        predicted_class = class_names[np.argmax(predictions[i])]
        true_class = class_names[sample_labels[i][0]]
        ax.set_title(f"Pred: {predicted_class}\nTrue: {true_class}")
        ax.axis('off')
    plt.tight_layout()
    plt.show()
```

```
[ ]: num_background = 100
background_images = test_images[:num_background]
explainer = shap.GradientExplainer(model, background_images)
```

```
[ ]: def shap_visualization(image_index):
    image = test_images[image_index:image_index+1]
    true_label = test_labels[image_index][0]

    # Generate and process SHAP values
    shap_values = explainer.shap_values(image)
    prediction = model.predict(image)
    predicted_class = np.argmax(prediction)
    shap_values_for_class = shap_values[0, ..., predicted_class]
    shap_sum = np.sum(shap_values_for_class, axis=-1)
```

```

# Normalize SHAP values for scatter plot
shap_normalized = (shap_sum - shap_sum.min()) / (shap_sum.max() - shap_sum.
↪min())

# Create figure with subplots
fig, axs = plt.subplots(1, 2, figsize=(12, 6))

# Original Image
axs[0].imshow(image[0])
axs[0].set_title("Original Image\nTrue: " + class_names[true_label])
axs[0].axis('off')

# Scatter Plot with Stars on Image
y, x = np.indices(shap_sum.shape)
colors = shap_sum.flatten() # Color by SHAP values
sizes = 100 * shap_normalized.flatten() + 10 # Size of stars
axs[1].imshow(image[0], aspect='auto') # Display the original image as ↪
↪background
scatter = axs[1].scatter(x.flatten(), y.flatten(), c=colors, s=sizes, ↪
↪cmap='coolwarm', marker='o', alpha=0.6)
axs[1].set_title("SHAP Scatter on Image\nPredicted: " + ↪
↪class_names[predicted_class])
axs[1].axis('off')
fig.colorbar(scatter, ax=axs[1], orientation='vertical', fraction=0.046, ↪
↪pad=0.04)

plt.tight_layout()
plt.show()

```

```
[ ]: shap_visualization(10)
```

1/1

0s 33ms/step

Original Image
True: airplane



SHAP Scatter on Image
Predicted: airplane

